

# Business Intelligence

Introduction to databases

# Lecture outline

- Database fundamentals
- Relational model and other data models
- Incomplete information
- Keys and constraints
- Query languages
- Examples

# Warning

- If you already took a database course, this lecture will be very easy for you

# Database fundamentals

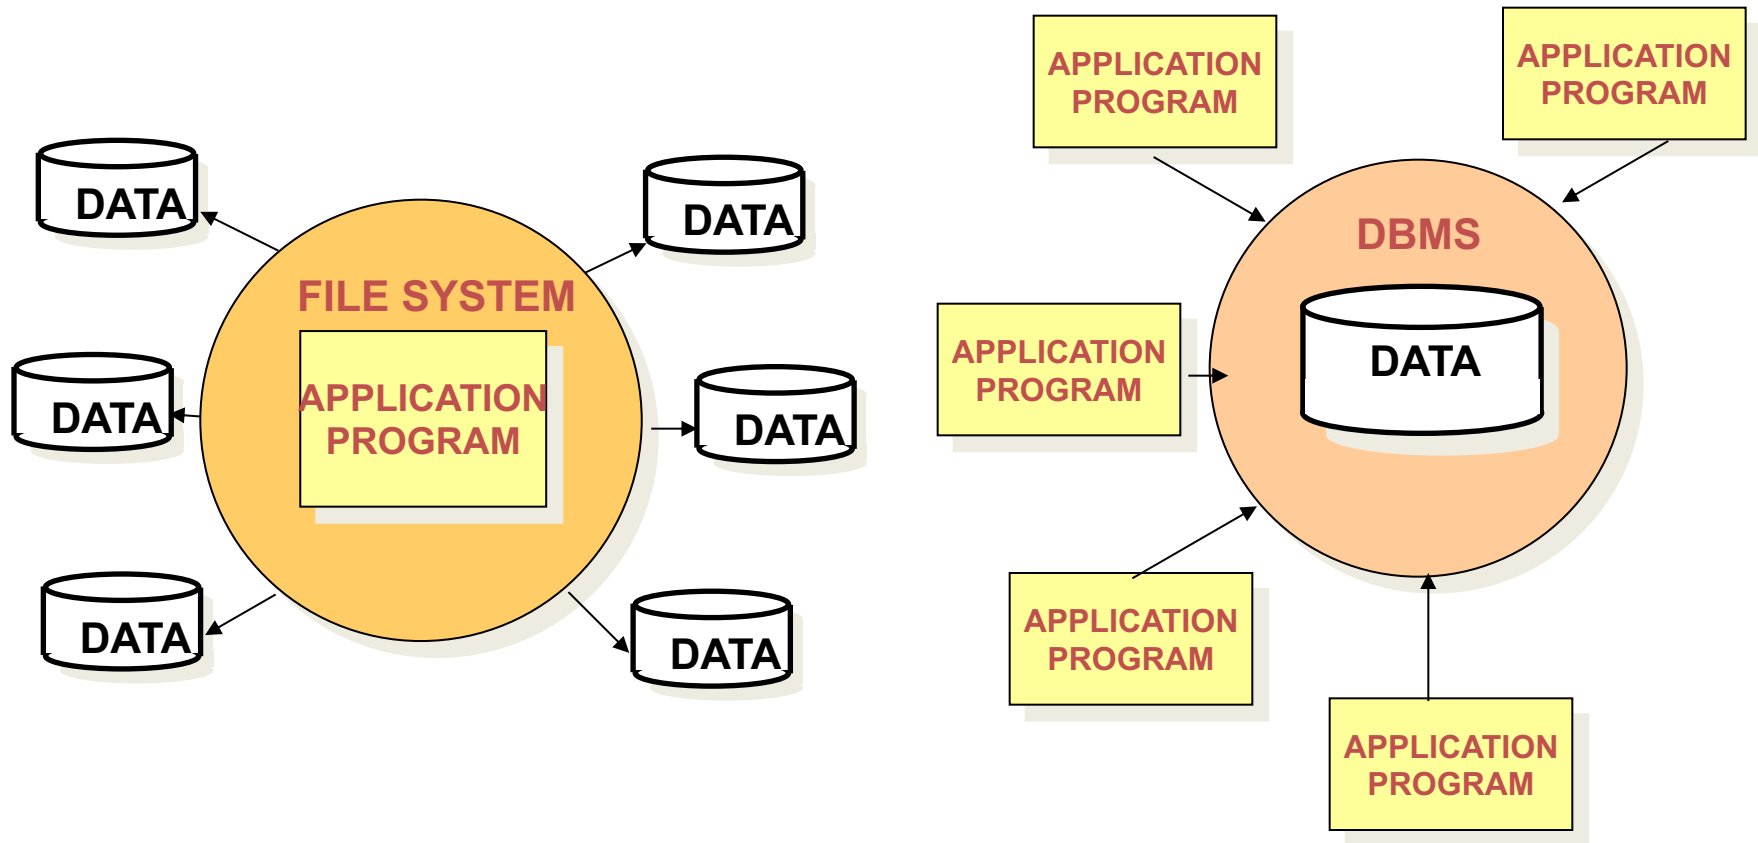
# Data and information

- **Data:**
  - elementary information unit
- **Information:**
  - processed data matching specific corporate needs

# Data and information

- Data:
  - < J. Smith,  
Business intelligence,  
spring semester >
- Information:
  - Who teaches BI? J. Smith
  - When does the course take place? In the spring semester

# Comparing databases and file systems



# Main features of DataBase Management Systems (DBMSs)

- Data sharing:
  - no replication in files
  - concurrency
- Data quality
  - integrity constraints
- Efficiency
  - loading, querying, sorting
- Access control
  - privacy
- Robustness




# Data integration

- All data, independently of the applications that use them, appear only once
  - (useless redundancies are introduced by design)
  - memory waste is reduced
  - data integrity is improved (the same data cannot have two different values at the same time)

# How to use a DBMS?

- Define the general structure of data
- Define the specific operations on data



# Example: university students

## Student

STUDNO	NAME	CITY	DEPTCODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

# The languages of a DBMS

- Data Definition Language (DDL)
  - examples: `CREATE`, `DROP`, `ALTER`
- Data Manipulation Language (DML)
  - examples: `SELECT`, `INSERT`, `UPDATE`, `DELETE`
- SQL (Structured Query Language) is a standardized language commonly used for both DDL and DML

# DML: Query Language

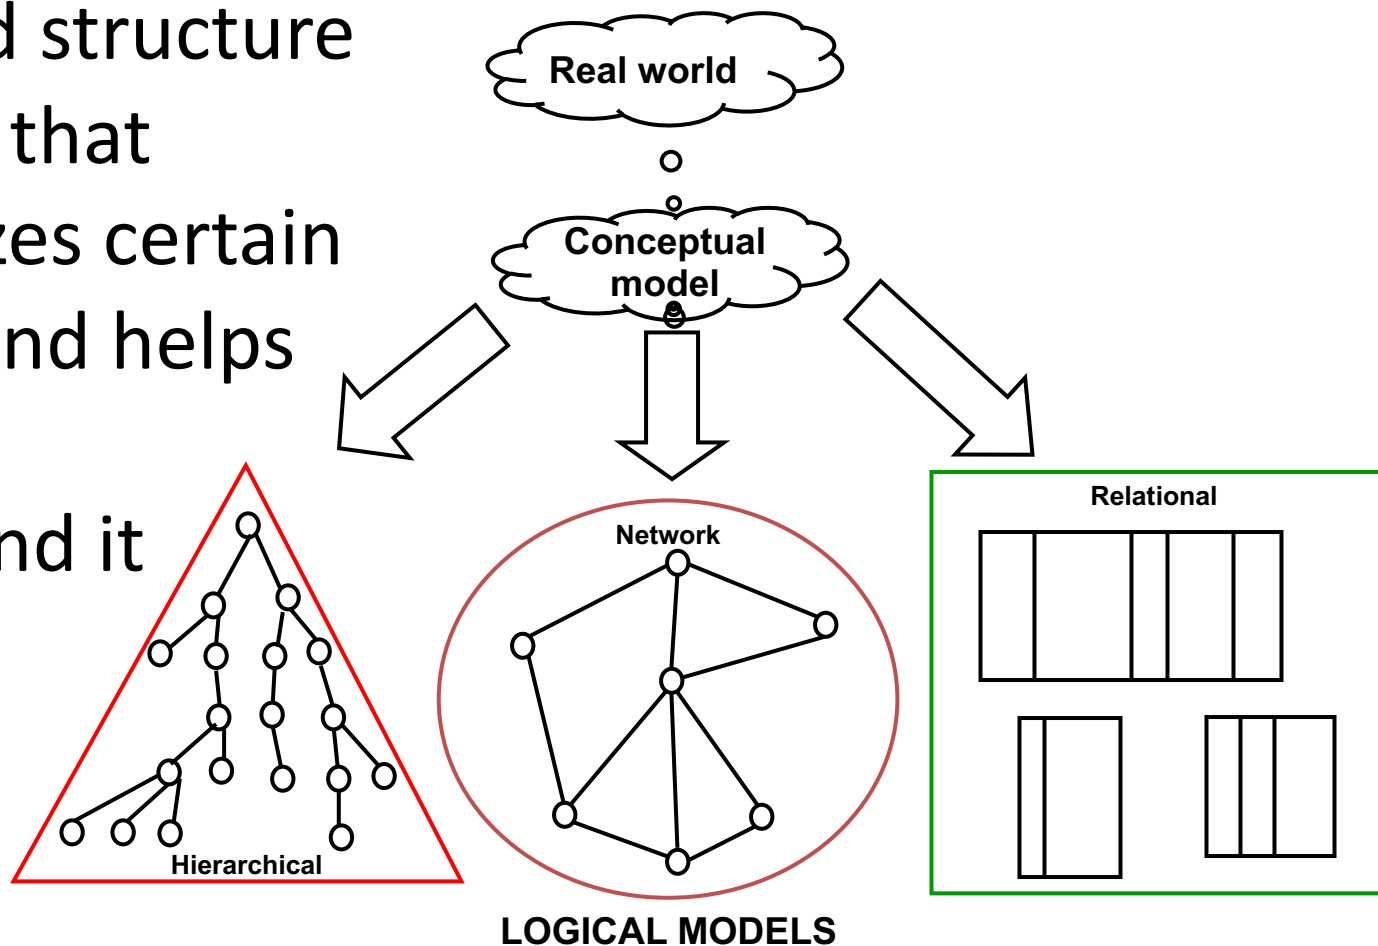
```
SELECT *  
FROM Student  
WHERE City = 'Lugano'
```

STUDNO	NAME	CITY	DEPTCODE
123	Jack	Lugano	Inf

# The relational model

# Data models

- Models offer a simplified structure of reality that emphasizes certain aspects and helps to better understand it




# Timeline of models for data representation

- Hierarchical model (1960s)
- Network model (1970s)
- Relational model (1980s)
- Object-oriented model (1990s)
- XML model (2000s)
- NoSQL model (2010s)




# Logical data models

- Hierarchical
  - Data are represented as records
  - Relationships between data are represented with **pointers** in a tree structure




# Logical data models

- Network (CODASYL consortium)
  - Data are represented as records
  - Relationships between data are represented with **pointers** in a complex graph structure



# Logical data models












































- Relational
  - Data are represented as tables
  - Relationships between data are obtained by associating attribute **values** in different tables



# History of the relational model

- Invented by Edgar F. Codd in 1969, published 1970 (IBM Research)
  - E. F. Codd: A Relational Model of Data for Large Shared Data Banks. Commun. ACM 13(6): 377-387 (1970)
- First projects: System R (IBM), Ingres (Berkeley University)
- Main technological findings: 1978-1980
- First commercial systems: start of 1980s (Oracle, IBM SQL/DS and DB2, Ingres, Informix, Sybase)
- Commercial success since 1985

# Popularity of the models

Rank			DBMS	Database Model	Score		
Sep 2023	Aug 2023	Sep 2022			Sep 2023	Aug 2023	Sep 2022
1.	1.	1.	Oracle 	Relational, Multi-model 	1240.88	-1.22	+2.62
2.	2.	2.	MySQL 	Relational, Multi-model 	1111.49	-18.97	-100.98
3.	3.	3.	Microsoft SQL Server 	Relational, Multi-model 	902.22	-18.60	-24.08
4.	4.	4.	PostgreSQL 	Relational, Multi-model 	620.75	+0.37	+0.29
5.	5.	5.	MongoDB 	Document, Multi-model 	439.42	+4.93	-50.21
6.	6.	6.	Redis 	Key-value, Multi-model 	163.68	+0.72	-17.79
7.	7.	7.	Elasticsearch	Search engine, Multi-model 	138.98	-0.94	-12.46
8.	8.	8.	IBM Db2	Relational, Multi-model 	136.72	-2.52	-14.67
9.	 10.	 10.	SQLite 	Relational	129.20	-0.72	-9.62
10.	 9.	 9.	Microsoft Access	Relational	128.56	-1.78	-11.47
11.	11.	 13.	Snowflake 	Relational	120.89	+0.27	+17.39
12.	12.	 11.	Cassandra 	Wide column, Multi-model 	110.06	+2.67	-9.06
13.	13.	 12.	MariaDB 	Relational, Multi-model 	100.45	+1.80	-9.70
14.	14.	14.	Splunk	Search engine	91.40	+2.42	-2.65
15.	 16.	 16.	Microsoft Azure SQL Database	Relational, Multi-model 	82.73	+3.22	-1.69
16.	 15.	 15.	Amazon DynamoDB 	Multi-model 	80.91	-2.64	-6.51
17.	 18.	 20.	Databricks	Multi-model 	75.18	+3.84	+19.56
18.	 17.	 17.	Hive	Relational	71.83	-1.52	-6.60
19.	19.	 18.	Teradata	Relational, Multi-model 	60.33	-0.98	-6.25
20.	20.	 24.	Google BigQuery 	Relational	56.46	+2.56	+6.34

# Informal definition

**column**

**schema**

**instance**

**row**

STUDNO	NAME	CITY	DEPTCODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco
234	Carl	Bern	Eco

# Relation and relationship

- In classical mathematics, a **relation** is a set-theoretic notion
- A **relationship**, a.k.a. association, indicates a correspondence between two entities in the model
  - We will see the Entity-Relationship model
- In the relational model, a relation has a slightly different meaning

# Formal definition: domain and Cartesian product

- A **domain**  $D$  is any set of values
- Let  $D_1, D_2, \dots, D_n$  be  $n$  (possibly not distinct) domains
- The **Cartesian product**

$$D_1 \times D_2 \times \dots \times D_n$$

is the set of all the ordered  $n$ -tuples

$$\langle d_1, d_2, \dots, d_n \rangle$$

such that  $d_1 \in D_1, d_2 \in D_2, \dots, d_n \in D_n$



# Example

- $D_1 = \{a, b\}$
- $D_2 = \{1, 2, 3\}$
- $D_1 \times D_2 = \{ \langle a, 1 \rangle, \langle b, 1 \rangle, \langle a, 2 \rangle, \langle b, 2 \rangle, \langle a, 3 \rangle, \langle b, 3 \rangle \}$

# Formal definition: mathematical relation

- A **mathematical relation** over  $D_1, D_2, \dots, D_n$  is a subset of the Cartesian product

$$D_1 \times D_2 \times \dots \times D_n$$

- $D_1, D_2, \dots, D_n$  are the **domains of the relation**
- A relation over  $n$  domains has **degree  $n$**
- The number of  $n$ -tuples is called the **cardinality** of the relation
- In real applications, the cardinality is always finite

# Example

- $D_1 = \{a, b\}$
- $D_2 = \{1, 2, 3\}$
- $D_1 \times D_2 = \{ \langle a, 1 \rangle, \langle b, 1 \rangle, \langle a, 2 \rangle, \langle b, 2 \rangle, \langle a, 3 \rangle, \langle b, 3 \rangle \}$
- $R_1 = \{ \langle a, 1 \rangle, \langle b, 3 \rangle \}$
- $R_2 = \{ \langle a, 1 \rangle, \langle b, 3 \rangle, \langle a, 2 \rangle \}$
- $R_3 = \emptyset$  (the empty set)
- $R_4 = \{ \langle a, 1 \rangle, \langle b, 1 \rangle, \langle a, 2 \rangle, \langle b, 2 \rangle, \langle a, 3 \rangle, \langle b, 3 \rangle \}$
- The degree of  $R_4$  is 2; its cardinality is 6
  - How many relations are there over  $D_1 \times D_2$ ?

# Properties

- Degree of a relation:
  - number of domains ( $n$ )
- Cardinality of a relation:
  - number of tuples
- Attribute:
  - name given to a domain in a relation
  - in a relation, the attribute names cannot be repeated

# Properties

- Based on the definitions, a mathematical relation is a set of ordered  $n$ -tuples
- Therefore, a relation is a set, and thus:
  - there is no ordering between the different  $n$ -tuples
  - the  $n$ -tuples of a relation are all distinct from one another
  - each single  $n$ -tuple is ordered: the  $i$ -th value of each tuple comes from the  $i$ -th domain, i.e., there is an ordering of domains

# Example

Games  $\subseteq$  string  $\times$  string  $\times$  integer  $\times$  integer

- each domain has two distinct roles, depending on their position (it is a **positional** structure):
  - the first and the third position regard the name and goals of the home team
  - the second and fourth regard name and goals of the visiting team

FC Basel	FC Lugano	3	1
FC Lugano	Servette FC	2	0
FC Basel	AC Bellinzona	1	2
AC Bellinzona	Servette FC	0	1

# Example under the relational model

- Each domain is associated with a name (**attribute**) describing its role in the relation
  - such a name is unique in the relation

Home	Visitor	HomeGoals	VisitorGoals
FC Basel	FC Lugano	3	1
FC Lugano	Servette FC	2	0
FC Basel	AC Bellinzona	1	2
AC Bellinzona	Servette FC	0	1

# Example under the relational model

- Attributes are also column headers in a table
  - The order of attributes is immaterial (non-positional structure)
  - These are the same:

Home	HomeGoals	VisitorGoals	Visitor
FC Basel	3	1	FC Lugano
FC Lugano	2	0	Servette FC
FC Basel	1	2	AC Bellinzona
AC Bellinzona	0	1	Servette FC

Home	Visitor	HomeGoals	VisitorGoals
FC Basel	FC Lugano	3	1
FC Lugano	Servette FC	2	0
FC Basel	AC Bellinzona	1	2
AC Bellinzona	Servette FC	0	1



# Formal definitions

- Let us associate each attribute with a domain, and let  $dom(A)$  indicate the domain associated with attribute  $A$
- An  $n$ -tuple over a set  $X$  of  $n$  attributes is a function that, for each attribute  $A$  in  $X$ , maps  $A$  to a value of the domain  $dom(A)$
- A **relation** over  $X$  is a set of  $n$ -tuples over  $X$

# Formal definition

- If  $t$  is an  $n$ -tuple over  $X$  and  $A \in X$ , then  $t[A]$  (or  $t.A$ ) indicates the value of  $t$  over  $A$
- In the example, if  $t$  is the first  $n$ -tuple of the table, then

$t[\text{Visitor}] = \text{'FC Lugano'}$

- The same notation is also extended to sets (actually, sequences) of attributes.

– It then denotes a tuple:

$t[\text{Visitor}, \text{VisitorGoals}] = \langle \text{'FC Lugano'}, 1 \rangle$

– (We do not care to distinguish between a value and a tuple with just one value)

# Tables and relations

- A table represents a relation if
  - The values of each column are from the same domain
  - Rows are different from one another
  - Column headers are different from one another
- Moreover, in a table representing a relation
  - the order of rows is irrelevant
  - the order of columns is irrelevant

# Comparing the terminology

Formal definition	Informal definition
relation	table
attribute	column
tuple, n-tuple	row
domain	data type
cardinality	number of rows
degree	number of columns

- An important difference:
  - formal definition: **no duplicates**
  - informal definition: duplicates are possible

# The relational model is value-based

- References between data in different relations are represented by means of values of the domains used in the  $n$ -tuples

Students	Studno	Last	First	DateOfBirth
	1234	Black	Joe	12/12/1990
	2345	White	John	11/11/1989
	3456	Red	Paul	10/10/1991
	4567	Green	Louise	08/08/1992

Exams

Student	Mark	Course
4567	A	01
4567	D	02
3456	B	04
1234	C	04

Courses

Code	Title	Teacher
01	Databases	Doe
02	Business Intelligence	Smith
04	Business Intelligence	Jones

Students

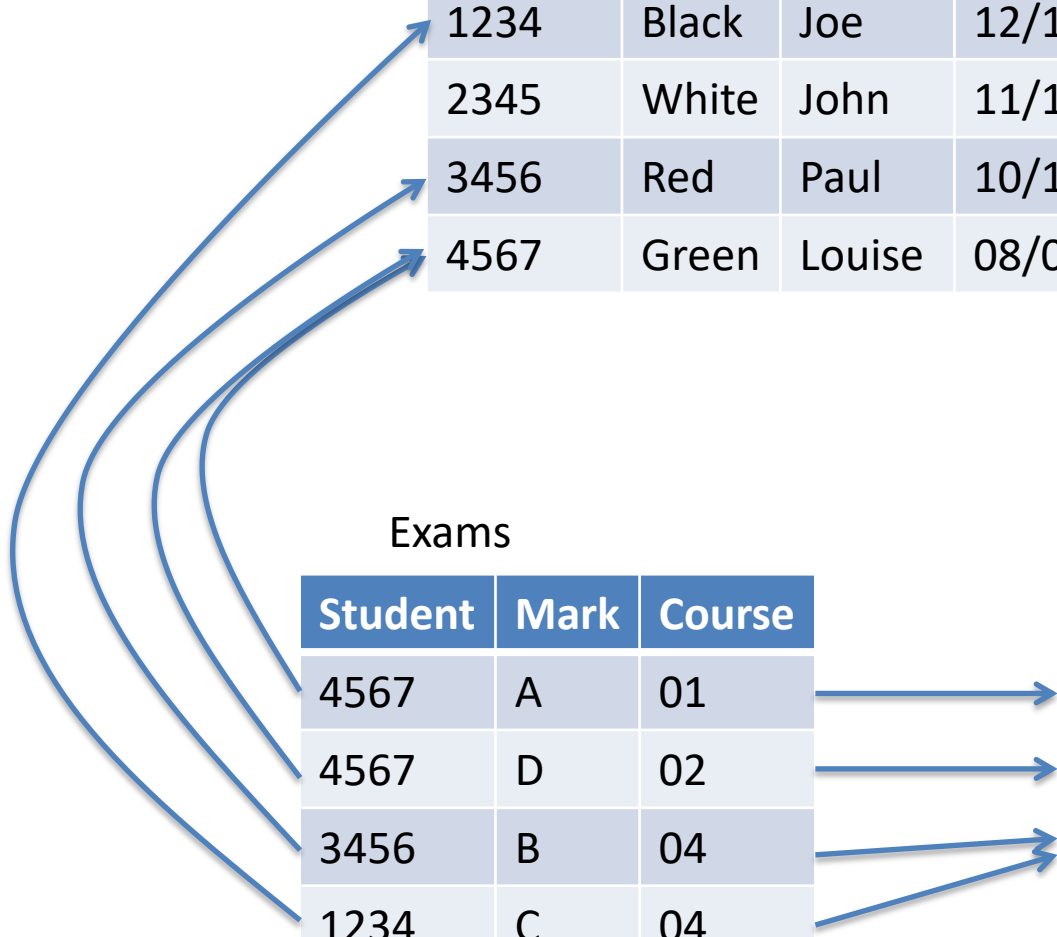
Studno	Last	First	DateOfBirth
1234	Black	Joe	12/12/1990
2345	White	John	11/11/1989
3456	Red	Paul	10/10/1991
4567	Green	Louise	08/08/1992

Exams

Student	Mark	Course
4567	A	01
4567	D	02
3456	B	04
1234	C	04

Courses

Code	Title	Teacher
01	Databases	Doe
02	Business Intelligence	Smith
04	Business Intelligence	Jones



# Why values and not pointers?

- Independence of physical structures
- Only what is relevant from the user application viewpoint is represented
  - pointers are less understandable for the end user
- Data are more easily portable to different systems
- Pointers are directional
  - and may exist at the physical level



# Cartesian product

Exam (E)

STUDNO	COURSECODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
702	2	7/9/13	C

Course (C)

COURSECODE	TITLE	TEACHER
1	BI	Doe
2	databases	Smith

E.STUDNO	E.COURSECODE	E.DATE	E.MARK	C.COURSECODE	C.TITLE	C.TEACHER
123	1	7/2/13	A	1	BI	Doe
123	2	8/1/13	B	1	BI	Doe
702	2	7/9/13	C	1	BI	Doe
123	1	7/2/13	A	2	databases	Smith
123	2	8/1/13	B	2	databases	Smith
702	2	7/9/13	C	2	databases	Smith

# Cartesian product

Exam

STUDNO	COURSECODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
702	2	7/9/13	C

Course

COURSECODE	TITLE	TEACHER
1	BI	Doe
2	databases	Smith

SELECT \* FROM Exam E, Course C

E.STUDNO	E.COURSECODE	E.DATE	E.MARK	C.COURSECODE	C.TITLE	C.TEACHER
123	1	7/2/13	A	1	BI	Doe
123	2	8/1/13	B	1	BI	Doe
702	2	7/9/13	C	1	BI	Doe
123	1	7/2/13	A	2	databases	Smith
123	2	8/1/13	B	2	databases	Smith
702	2	7/9/13	C	2	databases	Smith

# Queries

- Which professors have examined Jack?

Student

STUDNO	NAME	CITY	DEPTCODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

Exam

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
702	2	7/9/13	C

Course

COURSE CODE	TITLE	TEACHER
1	BI	Doe
2	databases	Smith

# SQL queries

- SQL queries have the typical `select-from-where` structure
- Syntax:  
`select AttrExpr {, AttrExpr}  
from Table {, Table}  
[ where Condition ]`
- The three parts of the query are called:
  - `select` clause with the target list
  - `from` clause
  - `where` clause
- The query
  - makes the “cartesian product” of the tables in the `from` clause
  - considers only the rows satisfying the condition in the `where` clause
  - for each row, evaluates the expression in the `select` clause
- Complete syntax:  
`select AttrExpr [[ as ] Alias ] {, AttrExpr [[ as ] Alias ] }  
from Table [[ as ] Alias ] {, Table [[ as ] Alias ] }  
[ where Condition ]`

# SQL query

- Which professors have examined Jack?

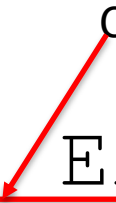
```
SELECT Teacher
FROM Course C, Student S, Exam E
WHERE Name = 'Jack'
AND S.StudNo = E.StudNo
AND E.CourseCode = C.CourseCode
```

# SQL query

- Which professors have examined Jack?

```
SELECT Teacher
FROM Course C, Student S, Exam E
WHERE Name = 'Jack'
AND S.StudNo = E.StudNo
AND E.CourseCode = C.CourseCode
```

Cartesian product



# Cartesian product

S.STUDNO	S.NAME	S.CITY	S.DEPTCODE	E.STUDNO	E.COURSECODE	E.DATE	E.MARK	C.COURSECODE	C.TITLE	C.TEACHER
123	Jack	Lugano	Inf	123	1	07/02/13	A	1	BI	Doe
456	Paula	Mendrisio	Inf	123	1	07/02/13	A	1	BI	Doe
789	Peter	Bellinzona	Eco	123	1	07/02/13	A	1	BI	Doe
123	Jack	Lugano	Inf	123	2	08/01/13	B	1	BI	Doe
456	Paula	Mendrisio	Inf	123	2	08/01/13	B	1	BI	Doe
789	Peter	Bellinzona	Eco	123	2	08/01/13	B	1	BI	Doe
123	Jack	Lugano	Inf	702	2	07/09/13	C	1	BI	Doe
456	Paula	Mendrisio	Inf	702	2	07/09/13	C	1	BI	Doe
789	Peter	Bellinzona	Eco	702	2	07/09/13	C	1	BI	Doe
123	Jack	Lugano	Inf	123	1	07/02/13	A	2	databases	Smith
456	Paula	Mendrisio	Inf	123	1	07/02/13	A	2	databases	Smith
789	Peter	Bellinzona	Eco	123	1	07/02/13	A	2	databases	Smith
123	Jack	Lugano	Inf	123	2	08/01/13	B	2	databases	Smith
456	Paula	Mendrisio	Inf	123	2	08/01/13	B	2	databases	Smith
789	Peter	Bellinzona	Eco	123	2	08/01/13	B	2	databases	Smith
123	Jack	Lugano	Inf	702	2	07/09/13	C	2	databases	Smith
456	Paula	Mendrisio	Inf	702	2	07/09/13	C	2	databases	Smith
789	Peter	Bellinzona	Eco	702	2	07/09/13	C	2	databases	Smith

# SQL query

- Which professors have examined Jack?

```
SELECT Teacher
FROM Course C, Student S, Exam E
WHERE Name = 'Jack'
AND S.StudNo = E.StudNo
AND E.CourseCode = C.CourseCode
```

Cartesian product

+  
equality of values  
across tables



# SQL query

- Which professors have examined Jack?

```
SELECT Teacher
FROM Course C, Student S, Exam E
WHERE Name = 'Jack'
AND S.StudNo = E.StudNo
AND E.CourseCode = C.CourseCode
```

Cartesian product

+

equality of values  
across tables

The two things together give rise to what is called a **join** in DB terms  
Joins have an alternative special syntax in SQL, which we won't see in this course

# Applying value equalities

S.STUDNO	S.NAME	S.CITY	S.DEPTCODE	E.STUDNO	E.COURSECODE	E.DATE	E.MARK	C.COURSECODE	C.TITLE	C.TEACHER
123	Jack	Lugano	Inf	123	1	07/02/13	A	1	BI	Doe
456	Paula	Mendrisio	Inf	123	1	07/02/13	A	1	BI	Doe
789	Peter	Bellinzona	Eco	123	1	07/02/13	A	1	BI	Doe
123	Jack	Lugano	Inf	123	2	08/01/13	B	1	BI	Doe
456	Paula	Mendrisio	Inf	123	2	08/01/13	B	1	BI	Doe
789	Peter	Bellinzona	Eco	123	2	08/01/13	B	1	BI	Doe
123	Jack	Lugano	Inf	702	2	07/09/13	C	1	BI	Doe
456	Paula	Mendrisio	Inf	702	2	07/09/13	C	1	BI	Doe
789	Peter	Bellinzona	Eco	702	2	07/09/13	C	1	BI	Doe
123	Jack	Lugano	Inf	123	1	07/02/13	A	2	databases	Smith
456	Paula	Mendrisio	Inf	123	1	07/02/13	A	2	databases	Smith
789	Peter	Bellinzona	Eco	123	1	07/02/13	A	2	databases	Smith
123	Jack	Lugano	Inf	123	2	08/01/13	B	2	databases	Smith
456	Paula	Mendrisio	Inf	123	2	08/01/13	B	2	databases	Smith
789	Peter	Bellinzona	Eco	123	2	08/01/13	B	2	databases	Smith
123	Jack	Lugano	Inf	702	2	07/09/13	C	2	databases	Smith
456	Paula	Mendrisio	Inf	702	2	07/09/13	C	2	databases	Smith
789	Peter	Bellinzona	Eco	702	2	07/09/13	C	2	databases	Smith

# Selecting columns

S.STUDNO	S.NAME	S.CITY	S.DEPTCODE	E.STUDNO	E.COURSECODE	E.DATE	E.MARK	C.COURSECODE	C.TITLE	C.TEACHER
123	Jack	Lugano	Inf	123	1	07/02/13	A	1	BI	Doe
456	Paula	Mendrisio	Inf	123	1	07/02/13	A	1	BI	Doe
789	Peter	Bellinzona	Eco	123	1	07/02/13	A	1	BI	Doe
123	Jack	Lugano	Inf	123	2	08/01/13	B	1	BI	Doe
456	Paula	Mendrisio	Inf	123	2	08/01/13	B	1	BI	Doe
789	Peter	Bellinzona	Eco	123	2	08/01/13	B	1	BI	Doe
123	Jack	Lugano	Inf	702	2	07/09/13	C	1	BI	Doe
456	Paula	Mendrisio	Inf	702	2	07/09/13	C	1	BI	Doe
789	Peter	Bellinzona	Eco	702	2	07/09/13	C	1	BI	Doe
123	Jack	Lugano	Inf	123	1	07/02/13	A	2	databases	Smith
456	Paula	Mendrisio	Inf	123	1	07/02/13	A	2	databases	Smith
789	Peter	Bellinzona	Eco	123	1	07/02/13	A	2	databases	Smith
123	Jack	Lugano	Inf	123	2	08/01/13	B	2	databases	Smith
456	Paula	Mendrisio	Inf	123	2	08/01/13	B	2	databases	Smith
789	Peter	Bellinzona	Eco	123	2	08/01/13	B	2	databases	Smith
123	Jack	Lugano	Inf	702	2	07/09/13	C	2	databases	Smith
456	Paula	Mendrisio	Inf	702	2	07/09/13	C	2	databases	Smith
789	Peter	Bellinzona	Eco	702	2	07/09/13	C	2	databases	Smith

# Final result

C.TEACHER
Doe
Smith

# Queries

- Which students got an A in BI?

Student

STUDNO	NAME	CITY	DEPTCODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

Exam

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
702	2	7/9/13	C

Course

COURSECODE	TITLE	TEACHER
1	BI	Doe
2	databases	Smith

# SQL query

- Which students got an A in BI?

```
SELECT Name
```

```
FROM Course C, Student S, Exam E
```

```
WHERE C.Title = 'BI'
```

```
AND E.Mark = 'A'
```

```
AND S.StudNo = E.StudNo
```

```
AND E.CourseCode = C.CourseCode
```

# Tables occurring more than once

- Table aliases are useful
  - For readability
  - For those cases where the same table occurs more than once
- Aliases are like variables in programming languages

# Who are Bo's employees?

Employee

EMPNO	NAME	SALARY	MGRNO
1	Al	100 K\$	2
2	Bo	200 K\$	NULL
3	Carl	150 K\$	2

```
SELECT X.Name  
FROM Employee X, Employee Y  
WHERE X.MgrNo = Y.EmpNo  
AND Y.Name = 'Bo'
```

X.Name
Al
Carl



# Selection and projection

- The operation indicated by the `SELECT` clause is called a **projection**
  - It restricts the tuples on those attributes indicated after the `SELECT` keyword
  - It keeps them all if `SELECT` is followed by `*`
- The condition indicated in the `WHERE` clause performs a so-called **selection**:
  - Only the rows satisfying the condition are retained
- It's a bit counterintuitive, but:
  - `SELECT` → projection
  - `WHERE` → selection
  - `FROM` → cartesian product

# Observations

- Difference between schema and instance
- Quite different activities:
  - schema design
  - instance management
- Moving from data to information (query language)

# Example

- Two instances of invoices

<b>"Chez Philippe"</b> <b>Via Elvezia 23</b> <b>9100 Somewhere</b>		
<b>Invoice no. 2369</b> <b>12/5/2012</b>		
<b>3</b>	<b>cover charge</b>	<b>3.15</b>
<b>2</b>	<b>starters</b>	<b>6.22</b>
<b>3</b>	<b>entrée</b>	<b>12.60</b>
<b>2</b>	<b>steaks</b>	<b>19.00</b>
<b>Total</b>		<b>41.98</b>

<b>"Chez Philippe"</b> <b>Via Elvezia 23</b> <b>9100 Somewhere</b>		
<b>Invoice no. 2456</b> <b>16/5/2012</b>		
<b>2</b>	<b>cover charge</b>	<b>2.10</b>
<b>1</b>	<b>starter</b>	<b>3.11</b>
<b>2</b>	<b>entrée</b>	<b>8.40</b>
<b>2</b>	<b>fish</b>	<b>25.5</b>
<b>2</b>	<b>coffee</b>	<b>1.60</b>
<b>Total</b>		<b><u>39.41</u></b>

# Relational representation, 1

## Invoice

Number	Date	Total
2369	12/5/2012	41.98
2456	16/5/2012	39.41

## Detail

Number	Quantity	Description	Amount
2369	3	cover charge	3.15
2369	2	starters	6.22
2369	3	entrée	12.60
2369	2	steak	19.00
2456	2	cover charge	2.10
2456	1	starters	3.11
2456	2	entrée	8.40
2456	2	fish	25.50
2456	2	coffee	1.60

# Relational representation, 2

Invoice

Number	Date	Total
2369	12/5/2012	41.98
2456	16/5/2012	39.41

Detail

Number	Row	Quantity	Description	Amount
2369	1	3	cover charge	3.15
2369	2	2	starters	6.22
2369	3	3	entrée	12.60
2369	4	2	steak	19.00
2456	1	2	cover charge	2.10
2456	2	1	starters	3.11
2456	3	2	entrée	8.40
2456	4	2	fish	25.50
2456	5	2	coffee	1.60

Incomplete information

# Incomplete information

- The relational model imposes a rigid data structure:
  - information represented by n-tuples
  - only some n-tuples formats are allowed:
    - those that match the relation schema
- Available data may not always match the required format exactly, for several reasons

# Incomplete information

Driver	License_no
Alice	A123456
Jim	
Bob	
Dave	

- Jim has a license but we do not know its number
- Bob does not have a license
- We do not know whether Dave has a license or not



# Incomplete information

- Albeit a rather common practice, it is better not to use values of the domain (such as 0, “”, etc.) for representing incomplete information:
  - maybe there are no “unused” values in the domain
  - some “unused” values might become used later
- Such values require special care each time they are encountered

# Incomplete information

- A rough but effective technique:
  - a **null value** (**NULL**) denotes the absence of a value of the domain (**NULL** is not part of the domain)
  - Formally, extending the notion of n-tuple is sufficient:
    - $t[A]$  is a value of  $\text{dom}(A)$  for every attribute  $A$ 
      - or it is the null value **NULL**
  - Restrictions on the presence of null values are needed

# Incomplete information

Student

STUDNO	NAME	CITY	DEPT-CODE
123	Jack	Lugano	NULL
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

Exam

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
NULL	2	8/1/13	NULL
702	2	7/9/13	C

Course

COURSE CODE	TITLE	TEACHER
1	NULL	Doe
2	databases	Smith

# NULL

- Three different cases of incomplete information:
  - value is **unknown**: a value of the domain applies, but we do not know which one
  - value is **non-existent**: no value of the domain applies
  - **no information**: it is not known whether a value of the domain applies or not
  - DBMSs do not distinguish between these kinds of null values, and therefore they implicitly adopt the “no information” semantics

# Keys and constraints

# Integrity constraints

- There are database instances that are syntactically correct but that do not represent feasible information for the application

Exams

Student	Grade_perc	Course
123	92	01
123	106	02
234	67	02
345	85	03

# Integrity constraints

- Aim: excluding those instances that do not correspond to meaningful information in the application being represented
  - constraints on null values
  - key constraints
  - referential integrity
  - generic constraints

# Integrity constraints

- Definition of integrity constraint
  - property that must always be kept satisfied by every instance of the relations
  - constraints can be regarded as logical formulas that map every database instance to either **true** or **false**



# Integrity constraints

- Types of constraints:
  - intra-relation constraints.
    - domain constraints (i.e., constraints on values)
      - constraints on null values
    - tuple constraints
    - ...
  - inter-relation constraints
    - referential integrity constraints
    - ...

# Integrity constraints

- Useful for describing the world of interest in a more accurate way than just by the schema structure
- Data quality
- Useful design tool
- May be used by the system for query optimization

# Tuple constraints

- A condition on the values of each single n-tuple, independently of the other n-tuples
- Possible syntax: Boolean expression (with AND, OR, NOT) of atoms comparing attribute values or arithmetic expressions thereof
- Example (course 1 has a higher required grade):  
(Grade\_perc >= 80) OR (course != 1)
- A **tuple constraint** is also a **domain constraint** if it only regards one attribute
  - Example:  
(Grade\_perc >= 60) AND (Grade\_perc <= 100)

# Keys

- A **key** of a schema is a subset of the schema attributes that is **unique** and **minimal**
- Unique: no two tuples have the same key value
- Minimal: by removing any attribute from the key, uniqueness is lost
- A subset of attributes that is unique (but not necessarily minimal) is called a **superkey**

# Keys in the example

(underlined in red)

Student

<u>STUDNO</u>	NAME	CITY	DEPT-CODE
---------------	------	------	-----------

Exam

<u>STUDNO</u>	<u>COURSECODE</u>	DATE	MARK
---------------	-------------------	------	------

Course

<u>COURSECODE</u>	TITLE	TEACHER
-------------------	-------	---------

# Schemas with multiple keys

- One of them is called **primary key**
- The other ones are the **alternate keys**

CLIENT

(CLIENT\_CODE, ADDRESS, SSN)

Primary key: CLIENT\_CODE

Alternate key: SSN

# Existence of keys

- Relations are sets, therefore each relation may contain the same tuple only once
  - the set of all attributes of a relation schema is always a superkey of the relation
- Since the set of attributes is finite, every relation schema always has (at least) one key

# Importance of keys

- Existence of keys guarantees accessibility of all data in the database
- Every single value is univocally accessible through:
  - the name of the relation
  - the value of the key
  - the name of the attribute
- Keys are the main means to connect data in different relations
  - “the relational model is value-based”



# Keys and null values

- In the presence of null values for the attributes forming the key
  - the identification of the corresponding n-tuple is not possible
  - references to other relations are also affected
- The presence of null values in keys must be limited
- Practical solution: for every relation we choose a primary key for which null values are not allowed

# Foreign keys

- Pieces of information in different relations are connected by means of common values
  - In particular, (primary) key values
- A **referential integrity constraint** between relations  $R_1$  and  $R_2$  over attributes  $X$  imposes that the values over  $X$  of every n-tuple of the instance of  $R_1$  also occur as values of the (primary) key of the instance of  $R_2$

# Foreign keys in the example

(green arrows)

Student


<u>STUDNO</u>	NAME	CITY	DEPT-CODE
---------------	------	------	-----------

Exam

<u>STUDNO</u>	<u>COURSECODE</u>	DATE	MARK
---------------	-------------------	------	------

Course

<u>COURSECODE</u>	TITLE	TEACHER
-------------------	-------	---------



# Table definition

- A table consists of
  - An ordered set of attributes
  - A (possibly empty) set of constraints
- SQL's `create table` command defines a relation's schema and creates an empty instance

```
create table Student
(   StudNo      char(6) primary key,
    Name        varchar(30) not null,
    City         varchar(20),
    Dept-code    char(3)  )
```

# Intra-relation constraints

- Constraints are conditions that must be satisfied by every instance of the database
- Intra-relation constraints regard a single relation (two cases: tuple level or table level)
  - `not null` (on a single attribute; tuple level)
  - `primary key`: defines the primary key (once per table; entails `not null`)
  - `unique`: allows the definition of alternate keys (table level)
  - `check`: can represent several kinds of constraint

# Referential integrity

- It's a hierarchical (parent-child) relationship between tables
- Some attributes of the child table are defined as a *foreign key*
- The values contained in the foreign key must always be present in the parent table

```
create table Exam
(  StudNo      char(6) not null,
   Coursecode  char(3) not null,
   Date        date,
   Mark        ENUM('A','B','C','D','E','F'),
   primary key (StudNo,Coursecode),
   foreign key (StudNo) references Student(StudNo),
   foreign key (Coursecode) references Course(Coursecode)
)
```

# An incorrect instance

Student

STUDNO	NAME	CITY	DEPT-CODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

Course

COURSECODE	TITLE	TEACHER
1	BI	Doe
2	databases	Smith

Exam

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
789	2	7/9/13	C
123	2	1/8/14	A
702	NULL	1/2/15	NULL
555	1	3/4/16	B

Violates primary key

Violates not null

Violates foreign key

# Managing orphans

Student

STUDNO	NAME	CITY	DEPT-CODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

Exam

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
789	2	7/9/13	C



# Managing orphans

Student

STUDNO	NAME	CITY	DEPT-CODE
123	Jack	Lugano	Inf
456	Paula	Mendrisio	Inf
789	Peter	Bellinzona	Eco

Exam

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
789	2	7/9/13	C

Orphan: tuple without parent  
because of deletions or updates  
in the parent table

# Reacting to change

- Deleting/updating a parent's tuple may cause a violation of referential integrity
- Possible reactions:
  - cascade: propagates the change
  - set null: the foreign key is set to null
  - set default: the reference is assigned a default value
  - no action: disallows the operation
- Syntax:

```
on < delete | update >  
    < cascade | set null | set default | no  
action >
```

# Reacting to change: deletion

What happens to exams if a student is deleted?

- `cascade`  
the exams of the deleted student are deleted, too
- `set null`  
the StudNo of the the exams of the deleted student are set to null
- `set default`  
the StudNo of the the exams of the deleted student are set to a default value
- `no action`  
student deletion is disallowed if there are exams for that StudNo

# Reacting to change: update

What happens to exams if a student's StudNo is changed?

- `cascade`  
the StudNo of the exams of the updated student are updated, too
- `set null`  
the StudNo of the the exams of the updated student are set to null
- `set default`  
the StudNo of the the exams of the updated student are set a default value
- `no action`  
student update is disallowed if there are exams for that StudNo

# Example syntax

```
create table Exam
( ....
    primary key (StudNo, CourseCode)
    foreign key (StudNo)
        references Student (StudNo)
        on delete cascade
        on update cascade
    foreign key (CourseCode)
        references Course (CourseCode)
        on delete no action
        on update no action )
```

# Ordering

- It's useful to sort results by relevance
- SQL has an `order by` clause
- Syntax:

`order by Attr [asc | desc] {, Attr [asc | desc]}`

- The sorting conditions are evaluated one after the other
  - If there is a tie on the value of the first `order by` attribute, the second one is considered, and so on

# Sorting the result

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
123	2	8/1/13	B
702	2	7/9/13	C
555	1	3/4/16	B

Select \* from exam  
order by mark, date desc

STUDNO	COURSE CODE	DATE	MARK
123	1	7/2/13	A
555	1	3/4/16	B
123	2	8/1/13	B
702	2	7/9/13	C

# Aggregates

- Sometimes it's useful to extract statistical information from a set of values
- The easiest form is given by aggregate functions, which apply to a group of rows
  - `count`      (cardinality)
  - `sum`
  - `max`
  - `min`
  - `avg`      (average)



# The count operator

- `count` gives the number of rows or of distinct values; syntax:

```
count(< * | [distinct | all] AttrList >)
```

- Extract the number of students:

```
select count(*)  
from Student
```

- Extract the number of distinct values of attribute Name among all rows of Student:

```
select count(distinct Name)  
from Student
```

# sum, max, min, avg

- **Syntax:**

`< sum | max | min | avg > ([ distinct | all ] AttrExpr )`

- The `distinct` option considers each value only once
  - Useful only for `sum` and `avg`
- The `all` option is the default and considers all values different from *null*

# A query with sum

```
select sum(amount) as tot  
from Order  
where ClientNo=2
```

Order

ORDNO	CLIENTNO	DATE	AMOUNT
1	1	7/2/12	50,000,000
3	2	8/1/13	12,000,000
5	2	7/9/13	1,500,000
4	3	8/8/16	8,000,000
6	3	9/9/17	1,500,000
2	4	1/1/18	5,500,000

Tot
13,500,000

# Grouping

- Aggregates can be applied to a subset of the rows of a table thanks to the `group by` clause
- The groups themselves may be filtered via the `having` clause
- Example: extract the sum of amounts of orders placed after 8/8/12 for every client with at least 2 orders

```
select ClientNo, sum(Amount)
from Order
where Date > 8/8/12
group by ClientNo
having count(*) >= 2
```

# Step 1: selection

- Evaluating the `where` clause

ORDNO	CLIENTNO	DATE	AMOUNT
1	1	7/2/12	50,000,000
3	2	8/1/13	12,000,000
5	2	7/9/13	1,500,000
4	3	8/8/16	8,000,000
6	3	9/9/17	1,500,000
2	4	1/1/18	5,500,000

## Step 2: grouping

- Evaluating the group by clause

ORDNO	CLIENTNO	DATE	AMOUNT
3	2	8/1/13	12,000,000
5	2	7/9/13	1,500,000
4	3	8/8/16	8,000,000
6	3	9/9/17	1,500,000
2	4	1/1/18	5,500,000

# Step 3: computing aggregates

- Computing `sum (amount)` and `count (*)` for each group

CLIENTNO	SUM(AMOUNT)	COUNT(*)
2	13,500,000	2
3	9,500,000	2
4	5,500,000	1

## Step 4: group extraction

- Evaluating `having count (*) >= 2`

CLIENTNO	SUM(AMOUNT)	COUNT(*)
2	13,500,000	2
3	9,500,000	2
4	5,500,000	1



# Step 5: generating result

- Evaluating `select` clause

CLIENTNO	SUM(AMOUNT)
2	13,500,000
3	9,500,000

# Coherence between group by and target list

- Incorrect:

```
select Mark
from Exam
group by StudNo
```

← Which mark? For which exam in the group?

- Incorrect:

```
select E.CourseCode, count(*), C.Teacher
from Exam E, Course C
where E.CourseCode= C.CourseCode
group by E.CourseCode
```

Here the teacher is univocally determined by the CourseCode (because of the foreign key), but this might not have been the case

- Correct:

```
select E.CourseCode, count(*), C.Teacher
from Exam E, Course C
where E.CourseCode= C.CourseCode
group by E.CourseCode, C.Teacher
```

# Multiple grouping

- Extract the sum of quantities of details of orders placed by each client on each product, provided that the sum is above 50

Order	ORDNO	CLIENTNO	DATE	AMOUNT
Detail	ORDNO	PRODNO	QTY	

```
select ClientNo, ProdNo, sum(Qty)
from Order as O, Detail as D
Where O.OrdNo = D.OrdNo
group by ClientNo, ProdNo
having sum(Qty) > 50
```

# A possible result after join and grouping

- Extract the sum of quantities of details of orders placed by each client, provided that the sum is above 50

O.ORDNO	CLIENTNO	D.ORDNO	PRODNO	QTY
3	1	3	1	30
4	1	4	1	20
3	1	3	2	30
5	1	5	2	10
3	2	3	1	60
1	3	1	1	40
2	3	2	1	30
6	3	6	1	25

1,1 group

1,2 group

2,1 group

3,1 group

# Final result

- Computing `sum(Qty)` for the groups and evaluating the `having` clause

CLIENTNO	PRODNO	SUM(QTY)
1	1	50
1	2	40
2	1	60
3	1	95